

DC4400 Project Report

Extension of the BBC iPlayer off product schedules system to process change notifications

Oliver Matthew Bowker (220263618)

February 21, 2024



Contents

1 Executive Summary	6
2 Introduction	7
3 Background	9
3.1 Organisation and Team	9
3.2 Original Architecture	10
3.2.1 Initial design solution	12
4 Research	14
4.1 Storage Solutions and Parallelism	14
4.1.1 Redis and Elasticache	14
4.1.2 How thread-safety can be achieved	16
4.2 Agile and CI/CD	17
4.2.1 Software Changeover Strategies	19
5 Work Done	21
5.1 Requirements and epic creation	21
5.2 Investigation and spike	23
5.3 Slicing and kick-off	26
5.4 Build software	29
5.4.1 Coldstarts	29
5.4.2 Delta/change lambda	29
5.4.3 Garbage Collector	29
5.4.4 Integration tests	29
5.5 Release	30
6 Outputs	32
7 Future Work	33
7.1 Notifications direct to partners	33
7.2 Parallelise the current system	35
8 Conclusion	38
9 References	39
10 Appendix	45
10.1 Appendix A - Partnerships Objectives	45
10.2 Appendix B - Full schedules design including future external notification work	46
10.3 Appendix C - Initial flow diagrams for how events would be processed	47
10.4 Appendix D - Full ways of working diagram	49
10.5 Appendix E - Full software spike document	50

List of Figures

1	Bar chart showing growth of smart TVs in UK homes (Statista, 2023) <i>Created with Line Graph Maker</i> (Line Graph Maker, 2023)	7
2	Image showing SpaceChimps place in the BBC (Bowker, 2023)	9
3	Image showing catalogue hierarchy and how they reference each other.	10
4	Image showing the initial architecture.	10
5	Activity diagram showing schedule generators logic.	11
6	Image showing the initial updated schedules design (Lloyd, 2023).	12
7	Diagram showing how a broadcast list of an episode can become incorrect/polluted.	14
8	Difference between concurrency and parallelism. https://books.google.ie/books/about/Introduction_to_Concurrency_in_Programmi.html?id=J5-ckoCgc3IC&redir_esc=y	15
9	How redis pipelines don't guarantee sequential execution (Eyng, 2019).	15
10	Example of a deadlock scenario (Apache Software Foundation, 2013).	16
11	Sequence and activity diagrams outlining logic of optimistic locking.	17
12	SpaceChimps kanban board.	17
13	General pipeline flow.	18
14	Timeline of different changeover strategies (Banerjee, 2017).	19
15	Initiation stage of our ways of working.	21
16	Investigation stage of our ways of working.	23
17	Simple activity diagram showing basic concurrency logic for updating schedules triggered by a catalogue update.	25
18	Kick-off stage of our ways of working.	26
19	Pre-work slicing of work, organised vertically into components.	26
20	Implementation and migration diagram components (The open group, 2016 and Jonkers et al, 2011).	27
21	Figure showing archimate implementation/migration diagram for project.	28
22	Build stage of our of working.	29
23	Release stage of our ways of working.	30
24	Diagram showing access to keyspaces.	30
25	Automatic rollback architecture on alarm error.	30
26	Sequence diagram showing how partners become out of date in new system.	33
27	Potential architecture to serve notifications to partners.	34
28	How an schedules can be kept up to date using changesets.	34
29	Option for parallelised architecture.	36
30	AWS pipe including lambda enrichment stage.	37
31	Image taken from a presentation given at a partnerships context setting event (BBC Partnerships, 2023).	45

32	Full diagram of design for schedules pipeline, including future notifications to partners work (Lloyd, 2023)	46
33	Flow diagrams for schedule events (Lloyd, 2023)	47
34	Flow diagram for catalogue/programme events (Lloyd, 2023)	48
35	Full ways of working flow diagram used by SpaceChimp.	49

List of Tables

1	Table from study showing difference in issues found between approaches (Fairbanks, Tharigonda, Eisty, 2023).	18
2	Example of dynamoDB table columns that could be used.	35

1 Executive Summary

2 Introduction

The BBC is now over 100 years old (BBC, 2022) and is well known for it's TV channels and radio stations that are broadcast over the airwaves (Pilnick, Baer, 1973, p.3) to peoples homes across the UK. However this old way of broadcasting, sending out airwaves on a certain frequency to an antenna, is becoming less popular in the modern age of the internet. A study done by Ofcom showed that people '*watched on average about 16% less broadcast TV between 2019 ... and 2022*', with viewing '*decreasing by 47%*' (Ofcom, 2023, p.7) between ages 16-24. In addition another study carried out by media analyst firm Ampere found that in 2021 37% of people claimed to watch no linear TV, this increased to 45% by 2023 (Ampere Analysis, 2023).

This fall correlates with the significant rise in internet enabled TVs in the home, with statista finding that '*In 2014 just 11 percent of households in the UK owned a Smart TV, whereas, in 2023, nearly 74 percent of households reported owning a Smart TV.*' (Statista, 2023).

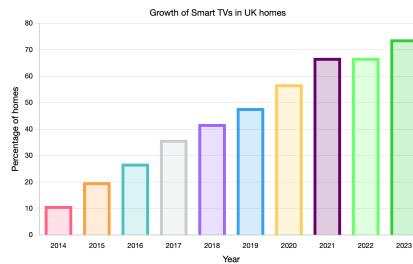


Figure 1: Bar chart showing growth of smart TVs in UK homes (Statista, 2023)
Created with Line Graph Maker (Line Graph Maker, 2023)

Some of these devices still support OTA broadcasts, however devices like the Amazon Fire TV stick and Googles Chromecast, are purely internet based; However they do offer a '*guide/epg*' section with Amazon having a development guide (Amazon, 2021) on how to integrate with it. Director general of the BBC, Tim Davie, in 2022 stated:

'The vision is simple: from today we are going to move decisively to a digital-first BBC' (Davie, 2022)

This statement highlights the goal to put more organisational focus on these new forms of media and internet enabled devices.

This report will discuss an upgrade carried out to the BBCs '*off-product*' schedules system, responsible for delivering up to date schedules to partners such as Freeview, Amazon and more. First I will give some background on the project, where I will discuss topics including storage Solutions and how they can work in parallel/multi-threaded systems, and strategies to protect live code

systems in a CI/CD environment. I will also give some background on the starting architecture of the system and how the changes align with the BBCs and teams OKRs (Sparks, 2024).

Following that, I will discuss the work that was done. This will be broken down into 5 sections that align with our teams ways of working flow.

1. Requirements and epic creation
2. Investigation and Spike
3. Slicing and task/ticket creation
4. Development of software
5. Releasing of software

I will then talk about the outputs of the project. These will include burn-up charts for the projects, dashboards created, documentation of the final architecture and a description of the final product.

Finally I will discuss potential improvements for future iterations. This will range between small code changes to a complete re-architecture of the system.

3 Background

In this section I will discuss the background work and research done for this project. I will start by discussing my teams place in the organisation and our OKRs, explaining how this project helps us hit these objectives. I will then outline the current architecture and the initial design for the project. Finally I will discuss some areas of interest around project, these include cloud computing, database parallelisation strategies and CI/CD challenges.

3.1 Organisation and Team

The BBC is broken into multiple layers with different responsibilities and goals. I am in a team called *SpaceChimps* which is part of the partnerships group, which itself is in the product group.

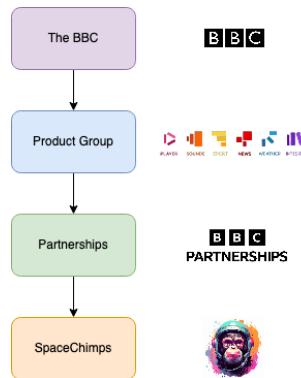


Figure 2: Image showing SpaceChimps place in the BBC (Bowker, 2023).

Our main aim as a team is to provide data to partners that they can use on their devices to promote BBC content. The project described in this report does just that by providing schedules for live content to partners. This aim fits directly into Partnerships objectives:

- The main objective of the schedules pipeline was to deprecate a service called *nitro*. This is an external platform the BBC used to distribute some of its data to partners. However this was a costly solution and meant we didn't have full control over the process. This projects aim is to improve this system allowing us to have better metrics and control over the data we share, providing the most up to date and relevant schedules.
- It helps drive growth as we are able to get content out to more people on more devices, increasing exposure to the BBC.
- It helps us improve our partner experience by working with them on integrating the data into their feeds.

- This project reduces the total time processing data, which therefore reduces our costs and makes us more sustainable.

All objectives can be seen in **Appendix A** (BBC Partnerships, 2023).

As well as schedules we also provide a '*catalogue*' of episodes, series and brands that are currently available on iPlayer.

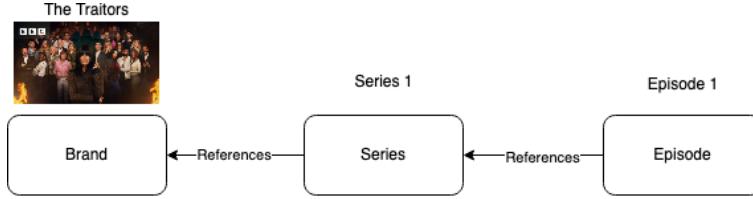


Figure 3: Image showing catalogue hierarchy and how they reference each other.

We provide both a 0 day and 8 day catalogue, these containing programme data that are available now or up to 8 days in the future. We also have an '*unfiltered*' catalogue that is not available to partners which contains all programme data with no availability limits. This unfiltered catalogue is what is used by the schedule pipeline to get it's data about episode/series/brands within the schedule.

3.2 Original Architecture

The original solution was composed of AWS services that created a pipe and filter architecture which transformed inputs into multiple outputs (Somerville, 2016, pp.182-183) that can be used by partners. The figure below shows this architecture.

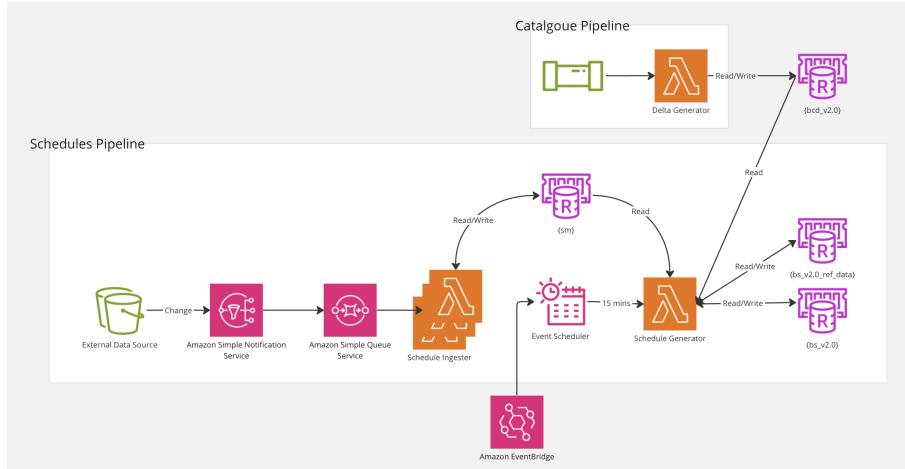


Figure 4: Image showing the initial architecture.

A pipe and filter architecture is achieved with the combination of AWS lambda (Amazon Web Services, 2024a) with updates being triggered by a publish/subscribe model (Somerville, 2021, pp.179) achieved through AWS' Simple Notification Service (SNS) (Amazon Web Services, 2024b) and Simple Queue Service (SQS) (Amazon Web Services, 2024c) with the former publishing and the latter subscribing. The lambda only runs when a message is published to the SQS, this triggers the lambda and the pipeline begins processing the new message.

The data is updated by an external system in AWS S3 (Amazon Web Services, 2024d), this publishes a message that the data has changed, our first lambda (the ingester) receives this message and processes/stores it into a '*common*' model that's used internally only. Our internal model is then up to date, however our partner facing model, created by the schedule generator, is not. In the original system this lambda was not driven by events, but instead ran every 15 minutes and processed all the schedules in one go, whether they had updated or not.

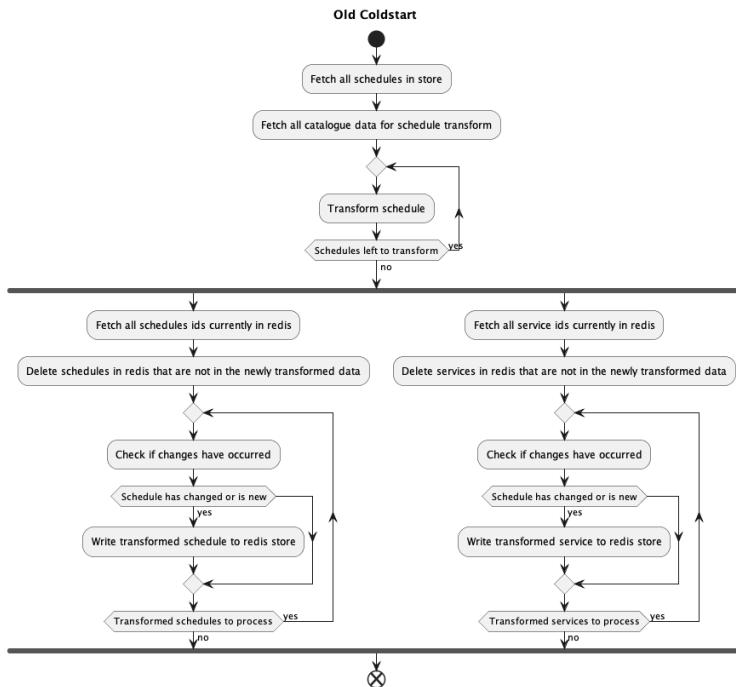


Figure 5: Activity diagram showing schedule generators logic.

This could leave partners with data up to 15 minutes out of date, dependant on when they retrieve the data, resulting in end users being shown incorrect schedules.

3.2.1 Initial design solution

The schedules pipeline also requires data from our catalogue pipeline, for titling, descriptions, viewer discretion warnings and subtitles. For this reason it needs to be alerted when catalogue data changes as well as when schedules change. An initial design had already been complete before the work started by another member of the team.

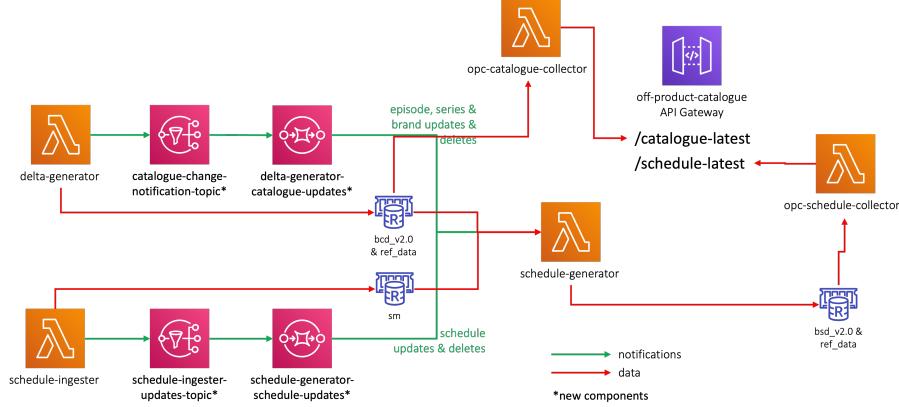


Figure 6: Image showing the initial updated schedules design (Lloyd, 2023).

The above diagram has been edited down to include only work related to the project, a full diagram including notifications being directly sent to partners can be found in **Appendix B**. This diagram also includes our endpoints and collectors that partners use to retrieve the data. This proposed system uses the pub/sub model described earlier and will subscribe to both catalogue and schedule events.

Alongside the architectural design an algorithm was proposed for these events.

- **For schedule updates** - We want to update the partner facing model of the schedule linked in the notification. This schedule has a list of broadcasts that map to episodes in the catalogue. A list should be maintained within each episode to create a link between both, allowing episode updates to trigger updates to schedules that reference them. All catalogue data referenced in a schedule should be copied over to the schedules keyspace, leaving the catalogue keyspace and items untouched.
- **For schedule deletes** - Remove schedule from store, and remove schedule reference from list contained within each episode that schedules references in its list of broadcasts.
- **For catalogue/programme updates:**
 - For episodes** - Update episode in store, update each schedule that is linked in episodes broadcast list.

For series/brands - Get all episodes linked to either the series or brand, update each schedule that is linked in each episodes broadcast list.

Flow diagrams can be found in **Appendix C** (Lloyd, 2023).

4 Research

I will now explore research done relating to the project specified. This will outline some design decisions that were made and link to other sections of the report where they may be explored more. This section will explore parallelism and thread safety focusing on data stores as well as CI/CD, it's pros and cons when deploying to live systems and how we as a team worked around it.

4.1 Storage Solutions and Parallelism

Currently the schedule pipeline consists of two components, the ingester, followed by the schedule generator. The first part of this pipeline is parallelised, multiple lambdas can be ran at the same time to insert data into the redis. This is a harder task for the schedule generator to do, as it needs update a list of linked schedules to an episode in the redis store. This data can be edited through multiple streams, both the schedule catalogue pipelines, meaning the array could easily become incorrect/polluted. Sharing memory in a threaded/parallelised system is a well known challenge and you need to know when it's safe to update/edit this memories value (<https://homes.cs.washington.edu/~dkg/teachingMaterials/spac/sophomoricParallelismAndConcurrency.pdf>). This is also known as being thread safe which can be described as '*different threads can access the same resources without exposing erroneous behavior or producing unpredictable results*' (Ugarte, 2024).

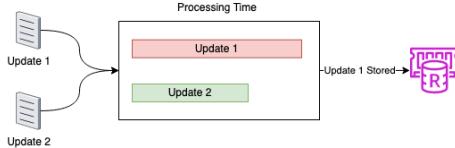


Figure 7: Diagram showing how a broadcast list of an episode can become incorrect/polluted.

4.1.1 Redis and Elasticache

We use amazons Elasticache (Amazon Web Services, 2024f) for our redis (REmote DIctionary Server) solution as it stores data in memory, which makes it really quick to retrieve stored data (IBM, 2024). This is vital for us as we store large documents that need to be retrieved and sent to partners on an API request. Redis is single-threaded but supports concurrency, '*when at least two threads are making progress*' (Oracle Corporation, 2010) which is not the same as parallelism, '*when at least two threads are executing simultaneously*' (Oracle Corporation, 2010).

Concepts in Concurrency

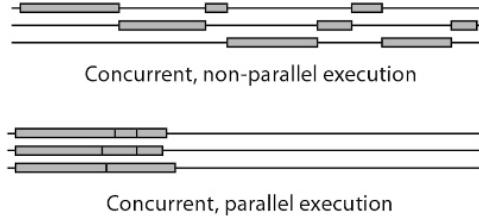


Figure 8: Difference between concurrency and parallelism.
https://books.google.ie/books/about/Introduction_to_Concurrency_in_Programmi.html?id=J5-ckoCgc3IC&redir_esc=y

This concurrency allows redis to support multiple requests at once, but it cannot do multiple operations at once; although more recent versions are allowing some safely threaded operations such as deleting records (Redis Ltd, 2024a). It also supports batch uploading and blocking commands, however these also don't help keep the data thread-safe.

- **Pipelining** - Pipelining sends a block of commands at once, however does not guarantee that commands sent are done in sequence (Redis Ltd, 2024b).

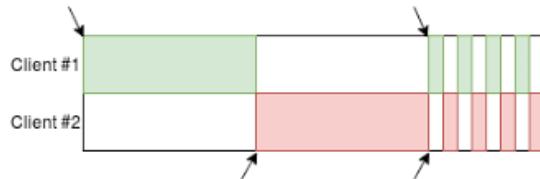


Figure 9: How redis pipelines don't guarantee sequential execution (Eyng, 2019).

- **Transactions** - Transactions are very similar to pipelines, however guarantee that the transactions commands are not interrupted by another clients requests and are therefore executed in sequence (Redis Ltd, 2024c).
- **Blocking Actions** - Blocking actions stop the current client from executing commands until the blocking action is complete. However other clients can still send requests to the server whilst this client is blocked (Redis Ltd, 2024d).

The options above still allow the previous race condition to occur, as instances of the schedule generator may vary in processing time and therefore the time of writing to the redis cannot be guaranteed to be in order of the events.

4.1.2 How thread-safety can be achieved

When researching and spiking (Visual Paradigm, 2024) the project, other technologies were found that could help offer thread safety whilst parallelising the schedule generator. These were types of store/database locking mechanisms.

- **Pessimistic Locking** - This method '*assumes that access to shared memory will be contended*' (Weston, 2011) and acquires a lock on the data to be edited. Any other client/connection attempting to edit this data must wait until this lock is released to update the data (Thornton, 2001). This can lead to issues such as deadlock which is when two clients are both awaiting on another clients lock to be released. This can end up with both clients being stuck in a endless cycle of waiting for each other (Thornton, 2001; Apache Software Foundation, 2013).

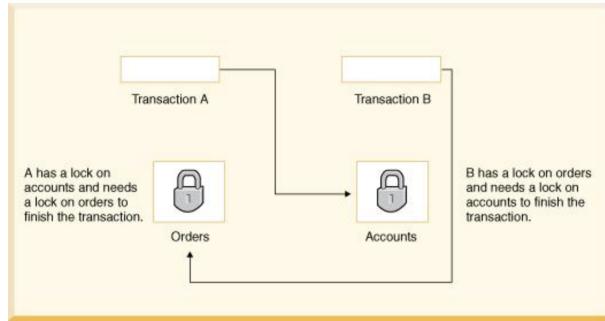


Figure 10: Example of a deadlock scenario (Apache Software Foundation, 2013).

There is also a situation where a client locks a piece of data and for one reason or another halts, this would cause needless delay, when a timeout is specified, if a timeout is not specified then this object would be locked from changes indefinitely (Thornton, 2001).

- **Optimistic Locking** - This method '*relies on end-of-transaction validation*' (Graefe, 2016) and takes the outlook of presuming it's safe to write until the very end (Kanungo, Morena, 2023). Unlike its counterpart (pessimistic locking) it does not lock the record that is being updated. Instead before the new data is written, the original data is checked against the current data stored (Thornton, 2001). If this data doesn't match then a change has occurred during processing and the new data to be written must be re-calculated with the new changes.

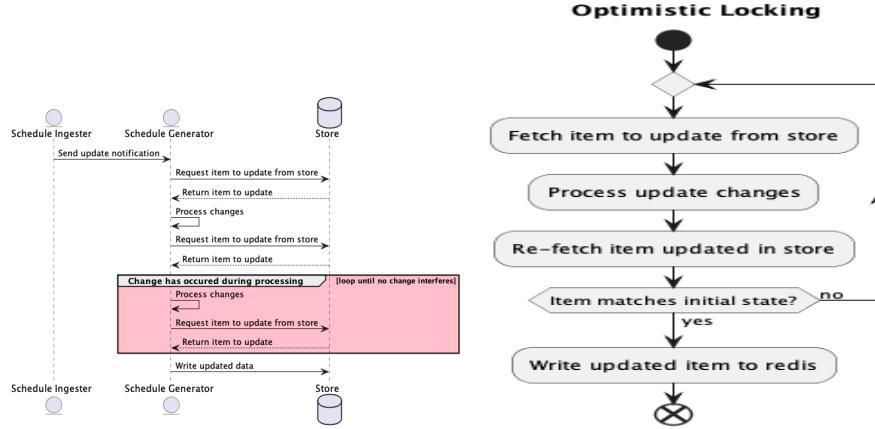


Figure 11: Sequence and activity diagrams outlining logic of optimistic locking.

This above logic could be implemented into our redis solution. It would require either a total comparison of the object or a simple *version* field that specifies when the object has changed.

During this investigation DynamoDB (Amazon Web Services, 2024e) was highlighted as a potential option due to it supporting optimistic locking. The final decision was to not use it however and stick with the current solution as there was a want to get the project started to keep up with our roadmap and not further investigate this option. There was also more unknowns with this technology as we had never used it before. I will discuss a solution using DynamoDB in the **Future Work** section of this report.

4.2 Agile and CI/CD

As a team we follow an agile approach to software development, more accurately we follow a kanban approach. The kanban methodology '*focuses more on monitoring and improving workflows*' (Heil, C, 2022) and doesn't use concepts such as sprints from scrum (Rehkopf, 2024). Instead software created using the kanban approach is deployed/released when it's done (Rehkopf, 2024) and uses a kanban board where tickets move across columns depicting where they are currently at in development cycle (Mauvius Group Inc, 2021).



Figure 12: SpaceChimps kanban board.

This workflow fits in well with Continuous Integrations and Continuous Deployment (CI/CD). As a team we release/build all our own code to multiple environments using Jenkins declarative pipelines (Jenkins, 2024).

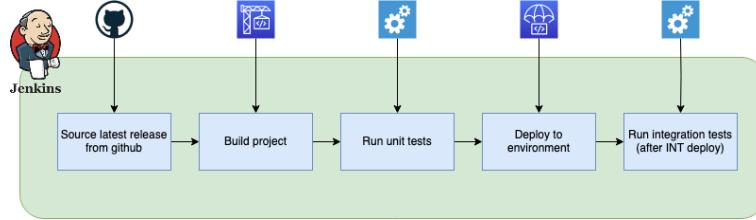


Figure 13: General pipeline flow.

Having these well defined deployment pipelines makes deploying new code easy and helps manage deployments across multiple environments as configuration can be built in to the pipeline (Rodriguez, et al, 2016). This approach allows bugs to be caught faster than when using a '*'big-bang'*' approach, as the bug is likely to be with a single new change (Department of Defence, 2021). The additional testing also gives more confidence in the code deployed and rollbacks are made extremely simple due to version control management.

There are of course some downsides to CI/CD, whenever an external platform, Jenkins in our case, is being used this opens up a new attack vector (NSA, 2023). OWASP keeps a list of some of these threats (OWASP Foundation, 2023), however they mostly comprise of poor authentication to CI/CD systems and attacks by people who already have access to source code and the CI/CD platform itself (employees).

Other issues include maintenance of the pipelines code/infrastructure and potential complexity (Wikström, 2019), this can be especially true when pipelines call other processes.

Code quality can become an issue, with technical debt increasing due to the encouraged continuous deployment of new software (Rodriguez, et al, 2016) and more bugs also appear to occur. One study showed that the number of bugs actually increased when using CI/CD (Fairbanks, Tharigonda, Eisty, 2023).

	CI/CD	No CI/CD
Github average issues	135.38	57.33
Gitlab average issues	52.04	17.68

Table 1: Table from study showing difference in issues found between approaches (Fairbanks, Tharigonda, Eisty, 2023).

However this same study also showed a commit velocity increase of 141%, so it could be argued that these bugs are quickly remedied due to this quicker release time. In addition to this as a team we do pull requests, another developer

checks new code before release, which should also mitigate some of the outlined issues above.

4.2.1 Software Changeover Strategies

The system that is being upgraded is in LIVE use by partners, but the work itself will take time to implement. We don't want to make changes to how the LIVE system works for partners but we do want to test that the new software works on the LIVE environment without a *big-bang* deployment.

There are three main types of changeover strategy, direct (big-bang), parallel running or a phased strategy and it's multiple variants (Banerjee, 2017).

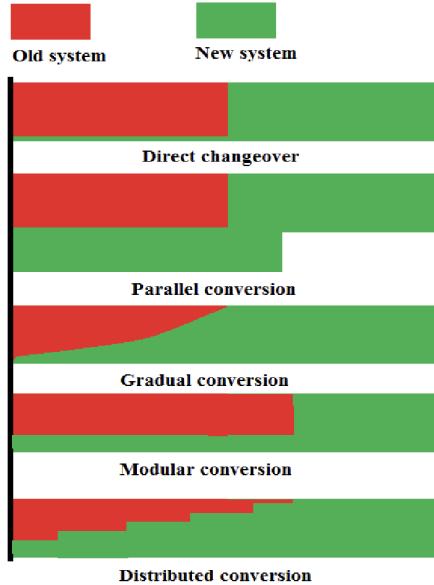


Figure 14: Timeline of different changeover strategies (Banerjee, 2017).

As discussed a direct approach is risky and could easily run into errors, a phased approach in this situation is also not valuable as the we want to remain consistent in the data we provide to all our partners, so a parallel system is our best option. Temporarily the old and new system will run side-by-side, this will allow us to write comparison tests between the old output and the new (Smyth, 2020).

In addition to having both systems the data on redis also needs to be separate as they may differ during development. This will be achieved by using different redis keyspace prefixes (IoRedis, 2024) to keep the data separate and allow a test to check both for differences (Rustagi, 2023). The partners would continue to get data from the old keyspace, when we are happy with our comparison tests the partners can be swapped over to the new keyspace and the old data can be

deleted.

5 Work Done

In this next section I will discuss the work that was carried on the project. In our team we have '*ways of working*' flowchart that helps guides us in how the project should be done throughout the team (GoRetro, 2023). Our ways of working is broken down into 5 sections, all of which will be discussed. For a full diagram of our workflow see **Appendix D**.

5.1 Requirements and epic creation

The First stage is initiation, what's the project going to be about, how it fits into the companies strategy and OKRs, requirement gathering and epic creation, an epic being a set of user stories/features (Karolis, Saulius, 2023).

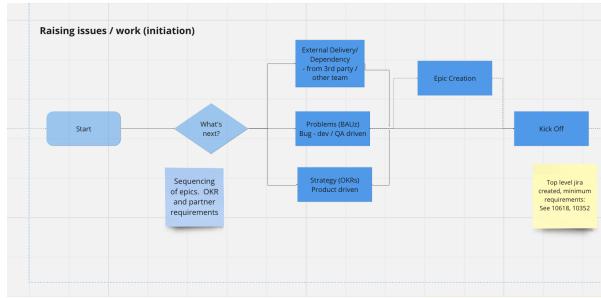


Figure 15: Initiation stage of our ways of working.

Due to complications in project allocation, I was not involved in the initial phases of development for this project. However discussions with people that were, highlighted that due to the lack of external partner interaction their wasn't much to discuss and no solid requirements were outlined, only the goal which is to allow the schedule generator to be moved from a static refresh of schedules to an event driven process.

The benefits of the project and how it achieves our OKRs has been discussed previously in the report. However I want to create some tangible requirements here using the MoSCoW framework which consists of must, should, could and won't haves (Eduardo, 2022). I will focus on software requirements and denote whether they are functional (F) or non-functional (NF) (Bigelow 2020).

1. (F) The software **Must** be able to process schedule updates.
2. (F) The software **Must** be able to process episode updates.
3. (F) The software **Must** be able to process ancestor updates.
4. (NF) The software **Must** keep stored data more relevant than the static version.
5. (F) The software **Must** still support a coldstart option.

6. (F) The software **Must** contain alarms to alert us to errors.
7. (F) The software **Should** gather metrics on processing time and types of updates.
8. (F) The software **Should** clean up unneeded data from it's store.
9. (NF) The software **Could** support parallelised processing.

In this case requirement gathering was easy as there were no partners to discuss the changes with as it's all internal and the output to partners stays the same. If this wasn't the case then requirements would have to be gathered from all partners affected or interested in the projects outcome. These requirements would be gathered through meetings and conversations with these partners individually and compromises may have to be made by both sides. This is usually the case when there is a requirement conflict, which can be described as:

'Requirements conflict is defined as unexpected or contradictory interaction between requirements that has a negative effect on the results (Cameron, Velthuijsen, cited in Kim et al, 2007)'

These conflicts need to be managed well, as Kim et al (2007) warns these conflicts can *'lead to negative or undesired operation of the system'*. However certain stakeholders will often carry more leverage than another, taking into consideration the Pareto Principle (Sanders, 1987), the 80-20 rule, it's easy for an organisation to leave out requirements of smaller stakeholders. A balance must be struck here where all parties are kept happy, without damaging the original idea of the project when large stakeholders try and morph it into what they deem it should be.

In this case no conflicts occurred due to the how internalised the project is. An epic was created in Jira and theses requirements could then be explored and implemented in the next phase.

5.2 Investigation and spike

After initiation the project is investigated and then *spiked*. In this case investigation/research for the most part has been done, as we knew the system was feasible due to following the same design as our catalogue pipeline. However it was determined that a spike of the proposed algorithm would be done.

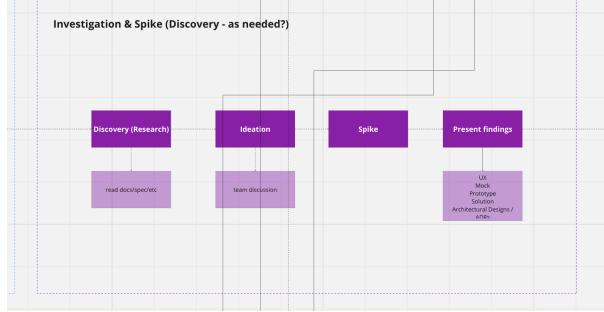


Figure 16: Investigation stage of our ways of working.

A spike is often used when there are unknowns about a project (Visual Paradigm, 2024). A study done by Hashimi and Gravell (2020) found that a majority of people in industry saw spikes as an effective tool and helped in managing risks. In another study they also stated that one of key purposes of a spike was to help guide story/ticket estimations (Hashimi, Gravell, 2019).

In the first paper it was also hypothesised that technical debt can also be lowered through the use of spikes, with technical debt be defined as:

'the idea that developers sometimes accept compromises in a system in one dimension (e.g., modularity) to meet an urgent demand in some other dimension (e.g., a deadline)' (Kruchten et al, 2012)

Spikes were used by Glas and Hedén (2021) to help bring down technical debt in their study. However when thought about simply, a spike could be seen as a tool to spot the technical debt before it touches live systems. Due to a prototype being made flaws can be spotted early and ironed out before the '*real*' development work begins.

For the project a spike was carried out to test the algorithms outlined in **Appendix C** to test the throughput of the system and see if it could handle the incoming requests. We used the technical spike template recommended by Microsoft (2021), to document findings and outline the spikes objectives. Hashimi, Abduldaem and Gravell (2022) found that '*timeboxing, objectives, documentation, and clear communication*' were the key factors that lead to a successful spike.

There was no time limit set, which did lead to scope creep (Martins, 2023) and more time being spent than was necessary. It's important to remember that a spike is not meant to be the final product, however scope creep lead to

the spike being almost a fully working system, albeit unrefined and with issues. This is something as a team we have now changed in our process and will always timebox spikes in the future to prevent this.

Despite this the other key factors were well adhered to. Objectives and documentation of the spike were put into a spike document and we had communication with a more senior member of the team at all times as well as daily stand ups to communicate issues/progress on the spike. The full spike document can be seen in **Appendix E**, below are the key findings.

1. A garbage collector should be used to clean up data that is no longer referenced by a schedule. This will help lower the amount of redis sets/gets and has no negative affect on partners.
2. When a catalogue item (episode/series/brand) updates it will also need to update it's associated schedules, for titling and descriptions. This in itself is not a problem, however if the number of schedules referenced is large it can significantly slow the system down. Parallelisation should be looked into, at least at the schedule level to help ease this. Full parallelisation would require a new design to support parallelised editing of the broadcast list held in episodes (this discussed in the **Research** section).
3. Redis duplication is complex and is only needed so that episodes can reference a list of schedules that they are in. Might be worth using DynamoDB here as this will also help with parallelisation (This is explored in the **Future Work** section).
4. Additional filtering from the catalogue pipeline could be added to only send updates for items that are referenced by a schedule. This would stop lambda being ran that essentially do nothing, this is also discussed in **Future Work**.

It was determined by the team that, for now, the system should be single-threaded and stick to the original design due to time constraints. Parallelisation could be looked into in the future when we had some time. However it was agreed for a garbage collector to be used and when a batch of schedule updates is triggered by a catalogue update that some form of concurrency would be required. This can be done with schedule objects due to the single threaded nature of the current system. As the schedule updates are being done because of a catalogue update the only fields that can change are the titles and descriptions. These do not affect or reference anything other than the object itself.

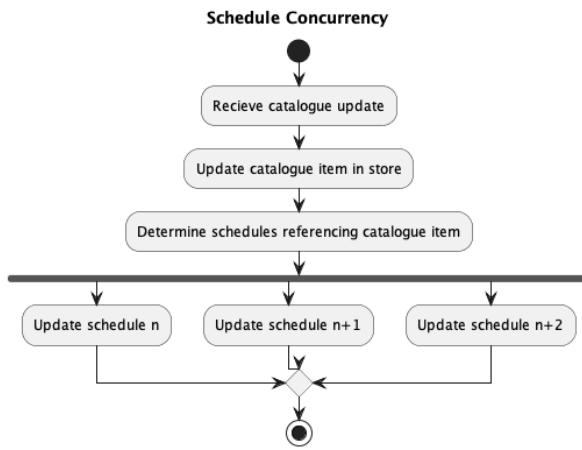


Figure 17: Simple activity diagram showing basic concurrency logic for updating schedules triggered by a catalogue update.

5.3 Slicing and kick-off

Another small section is the '*'kick-off'*', usually a single meeting where the team gets together and discusses the tickets that makeup the project as a whole.

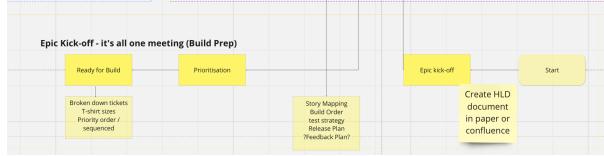


Figure 18: Kick-off stage of our ways of working.

Ideally this work has already been broken down into tasks. This is where the spike really helps, instead of guessing we are able to better understand the work that needs doing and create tasks accordingly (Hashimi, Abduldaem, Gravell, 2022). The following diagram shows the initial breakdown of tasks/tickets that were created for the project.

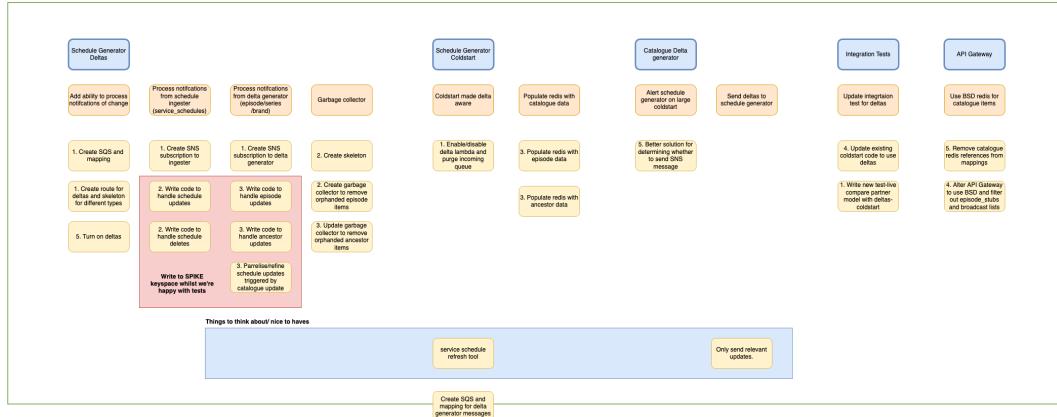


Figure 19: Pre-work slicing of work, organised vertically into components.

A work breakdown structure (WBS) is a tool used to '*'break down deliverables into sub-deliverables to visualize projects'* (Raeburn, 2024) and is typically broken down into 3 levels, parents, dependencies and sub-tasks, however it can consist of as a team wants. For the above I kept it at 3 levels, going vertically, blue represents self-contained features/components, orange is a parent task to complete the feature and yellow is a sub-task of the parent ticket. Numbers on the sub-tasks represent the initial prioritisation/ordering of the tasks into slices or feature sets.

It's important to note that this is not the final list of tasks. Unknowns will always appear whilst developing the software and result in additional tasks being created. However the WBS helps refine the scope to what is needed (Burghate,

2018). As can be seen on the above there are 3 tasks that have been moved to a '*nice to have*' section, which have been deemed unnecessary for the projects completion.

Another way to visualise these feature sets is using an Archimate implementation and migration diagram. This allows the modelling of work packages, gaps, deliverables, plateaus (stages) and events (The open group, 2016). The diagram on the next page shows the project mapped out using these elements, documenting the deliverables for each feature set.

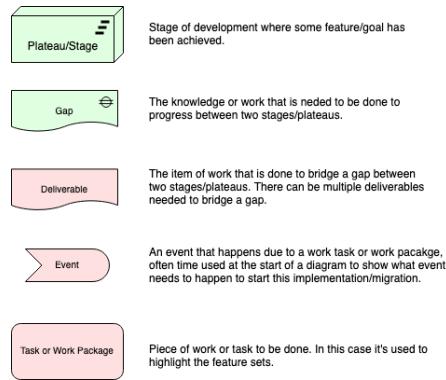


Figure 20: Implementation and migration diagram components (The open group, 2016 and Jonkers et al, 2011).

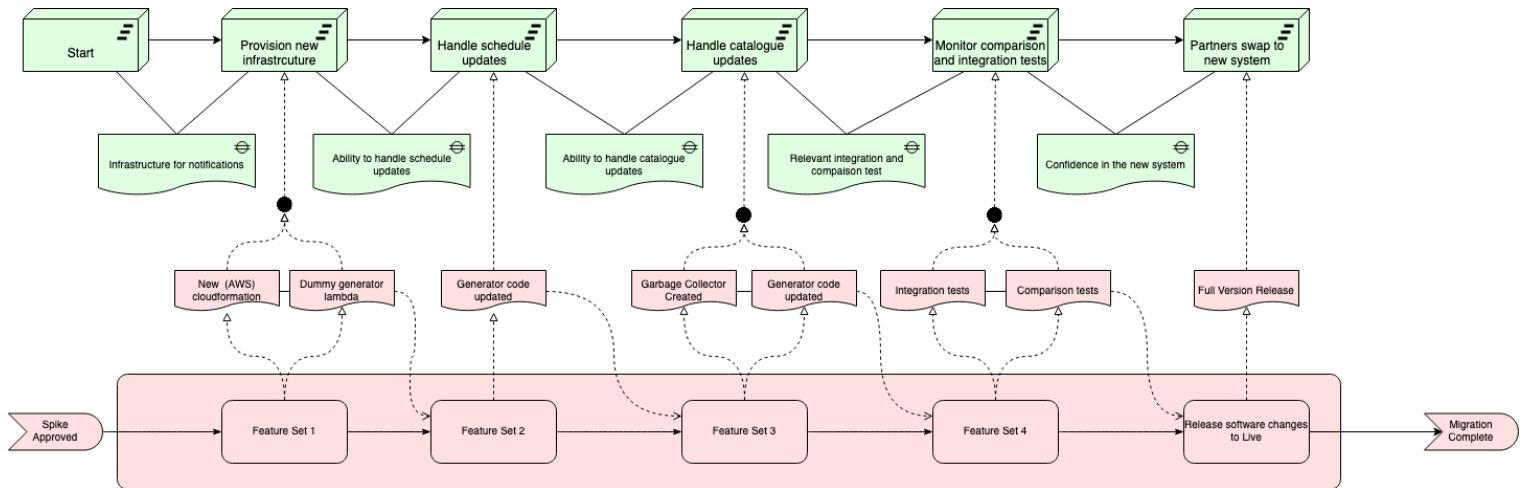


Figure 21: Figure showing archimate implementation/migration diagram for project.

5.4 Build software



Figure 22: Build stage of our of working.

5.4.1 Coldstarts

5.4.2 Delta/change lambda

5.4.3 Garbage Collector

5.4.4 Integration tests

5.5 Release

The release stage is the final stage in our workflow. This is when software gets released to the live environment and becomes available to partners.

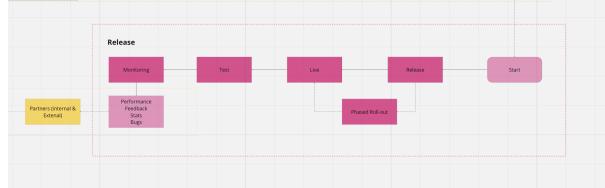


Figure 23: Release stage of our ways of working.

However in this case it doesn't become available to partners until we have confidence in the new system. We have config for each partner that specifies what data they have access to. In the case of schedules, this includes keyspace information on where the data is retrieved from. As previously discussed in the **Research** section, the new and old system will run in parallel until our comparison tests convince us that the new system outputs the same as the old.

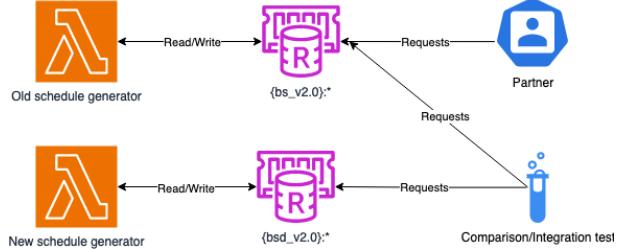


Figure 24: Diagram showing access to keyspaces.

Until then partners will still access data from the old system, with a simple config change being the rollover/rollback mechanism. If we wanted we could setup an automatic rollback and huddle system (Sathre, Zambreno, 2008), with the huddle portion consisting of developers being called out if outside of working hours to monitor and remedy the situation. This can be done in many ways but one way would be through AWS alarm actions (Amazon Web Services, 2024k). An alarm going into error would trigger a rollback of the configs version in S3 (Amazon Web Services, 2024d).

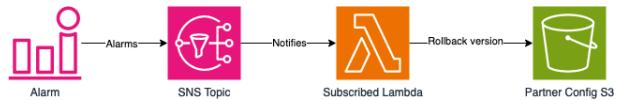


Figure 25: Automatic rollback architecture on alarm error.

For our case this is most likely over-engineering and unnecessary. A config change only takes a minute to do and if there ever was an error we would be notified and called out to fix the problem when necessary.

Once the switchover happens we wait for partner feedback and use created dashboards, these will be shown in the next section, to monitor the new software. Alongside this, our comparison test and side-by-side running of new and old systems will continue for a short while after release, allowing an easy rollback to the old system. Next this old lambdas scheduler will be disabled, but the lambda itself kept in case of an emergency. Finally, after sufficient time has passed the old system can be deconstructed and removed completely.

6 Outputs

7 Future Work

In this next section I will discuss two future upgrades that could be made to the schedules system to improve both its timeliness of data updates and a way that the system can be parallelised and simplified.

7.1 Notifications direct to partners

The new system relies on partners requesting data in order to update their UIs. This can mean that partners interfaces are still out of date, even though our internal store is as up to date as possible.

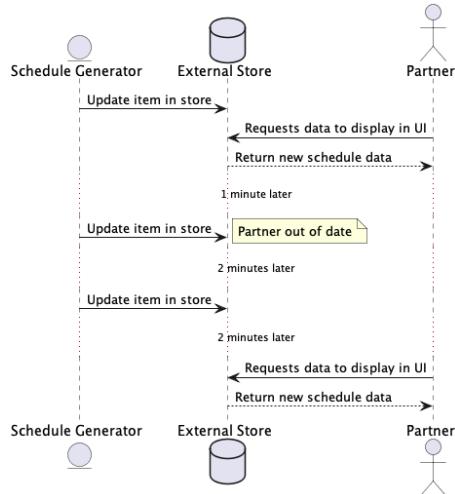


Figure 26: Sequence diagram showing how partners become out of date in new system.

In addition to this, re-requesting schedules that haven't been updated is a waste of time and costs the organisation money for data transfers (Amazon Web Services, 2024g). Below is an architecture that attempts to solve this problem by directly sending notifications to partners when an update occurs.

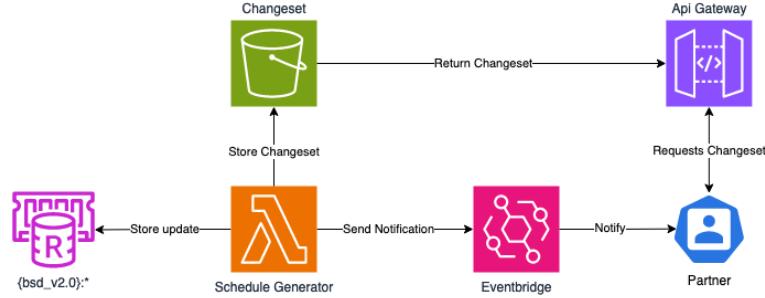


Figure 27: Potential architecture to serve notifications to partners.

A notification ID is sent to a partners system, they then request a '*changeset*' from our API which consists of only the data that has been updated. This new schedule can then be used instead of the old one, without doing a full refresh of all the schedules a client wants access to.

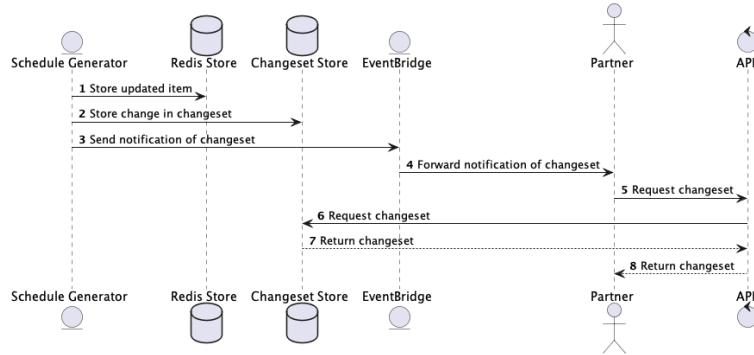


Figure 28: How schedules can be kept up to date using changesets.

Research

7.2 Parallelise the current system

The new system is single-threaded and is unable to run in parallel due to shared memory between the threads needing to be managed. The piece of memory at fault here is the broadcast list held in episodes that link that episode to schedules in which it is referenced. The main problem here is redis can't help us with any form of locking, unless we implement it ourselves, more on this can be found in the **Research** section of this report.

During this research DynamoDB came up as a possible solution due to it's use of optimistic locking capabilities (Kanungo, Morena, 2023 and Amazon Web Services 2024h). Using a column to monitor for changes, code could be written using Amazons Software Development Kit (SDK) (IBM Cloud Education, 2021) to ensure that no changes had occurred to the underlying before writing to the table.

ID	Associations	Version
Unique identifier of the object which would be used for quick lookups.	For series and brand objects this would be a list of children (series/episode) that the object refers to. For episodes this would now be where the broadcast list of related schedules is stored.	Number that is used to determine if changes have occurred using optimistic locking. This will be incremented by 1 every time a change happens.

Table 2: Example of dynamoDB table columns that could be used.

This table can be combined with the following SDK command to optimistically lock the data being edited. In this code *ConditionExpression* is the argument that determines what row is locking the record.

```
1 const updateCommand = new UpdateItemCommand({
2   TableName: 'schedule-associations',
3   Key: marshall({ id: objectIdToUpdate }),
4   UpdateExpression: 'set associations :newAssociations',
5   ConditionExpression: 'version = :previousVersion',
6   ExpressionAttributeValues: marshall({
7     ':newAssociations': newAssociations,
8     ':previousVersion': initialObject.version
9   })
10 })
```

Listing 1: SDK command sent to optimistically lock writes to the associations column.

Finally this can all be integrated into a new architecture that is outlined below.

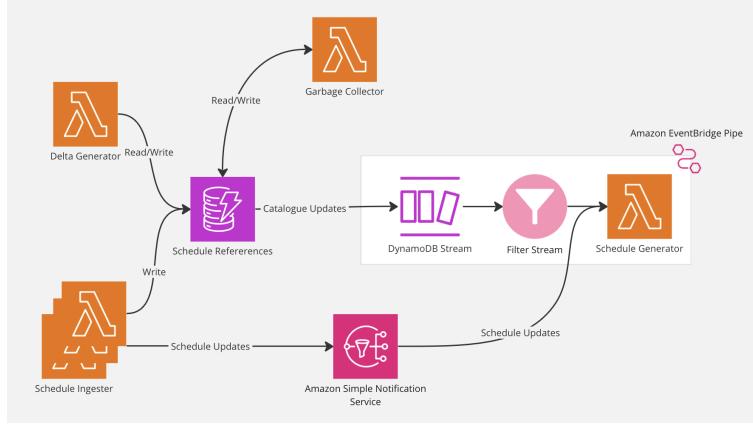


Figure 29: Option for parallelised architecture.

The system still takes input from both the schedule ingester and catalogue pipeline as before. However the aforementioned dynamoDB table is populated by ingestor and the catalogue pipeline, with the schedule generator no longer needing access to that data. Instead an AWS pipe (Amazon Web Services, 2024i) is used to trigger invocation of the delta lambda for catalogue updates, whilst schedule updates function the same. AWS pipes consist of 4 sections:

1. **Source** - DynamoDB streams can be used as a source to a pipe (Amazon Web Services, 2024j), this stream representing a change to the table. This event can contain both old and new items from the dynamoDB table, as well as the type of event that it is (Silveria, 2020).
2. **Filter** - We can filter out events based on the type of change made, we only want to trigger a lambda invocation on a MODIFY and not on an INSERT or DELETE as these events are covered by the standard schedule processing.
3. **Enrichment** - Not used in this example but can trigger extra enrichment of the data prior to the final target.
4. **Target** - This is where to send the filtered data to. In this case it's the schedule generator itself.

The inclusion of the dynamoDB table allows the use of optimistic locking to eventually parallelise the entire schedule pipeline. However there are other huge benefits here as well with the highlights being:

- There is no longer a need to copy over catalogue data to the schedules redis store as it was only used to maintain the link between episode and schedules via the broadcast list. This lowers the complexity of the schedule generator and allows the possibility for parallelism.

- The catalogue pipeline can be made to only send updates for objects that exist in schedules.

```

1      const updateCommand = new UpdateItemCommand({
2          TableName: 'schedule-associations',
3          Key: marshall({ id: updatedCatalogueItem.id }),
4          UpdateExpression: 'set version ${initialObject.
5              version + 1}',
6          ConditionExpression: 'id = ${updatedCatalogueItem.id
7              }',
8      })

```

Listing 2: SDK command sent by catalogue pipeline to ignore non schedule related catalogue items.

The above code once again uses the *ConditionExpression*, but instead of optimistically locking here it is used to check if the objects exists. This works as if the object doesn't exists in the table it can't have an id. This will stop needless invocations of the schedule generator.

- An additional refinement stage could be added to the pipe to check that the modification to data that has happened would cause a change to a schedule. Only certain fields changes in catalogue items cause this so this could again save more invocations.



Figure 30: AWS pipe including lambda enrichment stage.

8 Conclusion

9 References

- BBC. (2022) *Timeline - History of the BBC* [online]. Available at <https://www.bbc.co.uk/historyofthebbc/bbc-100/timeline/> (accessed on 18th January 2024).
- Pilnick, C. Baer, S. W. (1973) *Cable television: A guide to the technology* [online]. Available at <https://www.rand.org/content/dam/rand/pubs/reports/2006/R1141.pdf> (accessed on 18th January 2024).
- Ofcom. (2023) *Media Nations report 2023* [online]. Available at https://www.ofcom.org.uk/__data/assets/pdf_file/0029/265376/media-nations-report-2023.pdf (accessed on 18th January 2024).
- Ampere Analysis. (2023) *Almost half of Internet users say they have switched off broadcast TV* [online]. Available at <https://www.ampereanalysis.com/press/release/dl/almost-half-of-internet-users-say-they-have-switched-off-broadcast-tv> (accessed on 18th January 2024).
- Statista. (2023) *Household penetration of smart TV sets in the United Kingdom (UK) from 2014 to 2023* [online]. Available at <https://www.statista.com/statistics/654074/smart-tv-penetration-in-uk-households/> (accessed on 18th January 2024).
- Line Graph Maker. (2023) *Bar Graph Maker* [online]. Available at <https://linegraphmaker.co/bar-graph> (accessed on 19th January 2024).
- Amazon. (2021) *Linear TV Integration Guide Overview* [online]. Available at <https://developer.amazon.com/docs/fire-tv/linear-tv-integration-guide-overview.html> (accessed on 18th January 2024).
- Davie, T. (2022) *A digital-first BBC* [online]. Available at <https://www.bbc.co.uk/mediacentre/speeches/2022/digital-first-bbc-director-general-tim-davie> (accessed on 18th January 2024).
- Sparks, R. (2024) *OKRs: the ultimate guide to objectives and key results* [online]. Available at <https://www.atlassian.com/agile/agile-at-scale/okr> (accessed on 18th January 2024).
- Bowker O. (2023) *Analysis of BBC process and proposal for a new API key management system.* Assignment for DC4000, MSc Digital and Technology Solutions Specialist, Aston University, Unpublished.
- BBC Partnerships. (2023) *Partnerships - one year on.* BBC. Unpublished.
- Somerville, I. (2016) *Software engineering*. Tenth edition, Global edition. London: Pearson Education.
- Somerville, I. (2021) *Engineering Software Products: An Introduction to Modern Software Engineering*. Global edition. London: Pearson Education.
- Amazon Web Services. (2024a) *AWS Lambda* [online]. Available at <https://aws.amazon.com/lambda/> (accessed on 2nd February 2024).

Amazon Web Services. (2024b) *Amazon Simple Notification Service* [online]. Available at <https://aws.amazon.com/sns/> (accessed on 2nd February 2024).

Amazon Web Services. (2024c) *Amazon Simple Queue Service* [online]. Available at <https://aws.amazon.com/sqs/> (accessed on 2nd February 2024).

Amazon Web Services. (2024d) *Amazon S3* [online]. Available at <https://aws.amazon.com/s3/> (accessed on 2nd February 2024).

Amazon Web Services. (2024e) *Amazon DynamoDB* [online]. Available at <https://aws.amazon.com/dynamodb/> (accessed on 2nd February 2024).

Amazon Web Services. (2024f) *Amazon ElastiCache* [online]. Available at <https://aws.amazon.com/elasticache/> (accessed on 9th February 2024).

Amazon Web Services. (2024g) *How AWS Pricing Works* [online]. Available at <https://docs.aws.amazon.com/pdfs/whitepapers/latest/how-aws-pricing-works/how-aws-pricing-works.pdf> (accessed on 20th February 2024).

Amazon Web Services. (2024h) *Optimistic locking with version number* [online]. Available at <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/DynamoDBMapper.OptimisticLocking.html> (accessed on 21st February 2024).

Amazon Web Services. (2024i) *Amazon EventBridge Pipes* [online]. Available at <https://aws.amazon.com/eventbridge/pipes/> (accessed on 20th February 2024).

Amazon Web Services. (2024j) *Amazon DynamoDB stream as a source* [online]. Available at <https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-pipes-dynamodb.html> (accessed on 21st February 2024).

Amazon Web Services. (2024k) *PutMetricAlarm* [online]. Available at https://docs.aws.amazon.com/AmazonCloudWatch/latest/APIReference/API_PutMetricAlarm.html#API_PutMetricAlarm_Errors (accessed on 19th February 2024).

IBM. (2024) *What is Redis?* [online]. Available at <https://www.ibm.com/topics/redis> (accessed on 9th February 2024).

IBM Cloud Education. (2021) *SDK vs. API: What's the Difference?* [online]. Available at <https://www.ibm.com/blog/sdk-vs-api/> (accessed on 21st February 2024).

Redis Ltd. (2024a) *Redis FAQ* [online]. Available at <https://redis.io/docs/get-started/faq/> (accessed on 9th February 2024).

Redis Ltd. (2024b) *Redis pipelining* [online]. Available at <https://redis.io/docs/manual/pipelining/> (accessed on 9th February 2024).

Redis Ltd. (2024c) *Transactions* [online]. Available at <https://redis.io/docs/interact/transactions/> (accessed on 9th February 2024).

- Redis Ltd. (2024d) *Redis modules and blocking commands* [online]. Available at <https://redis.io/docs/reference/modules/modules-blocking-ops/> (accessed on 9th February 2024).
- Oracle Corporation. (2010) *Defining Multithreading Terms* [online]. Available at <https://docs.oracle.com/cd/E19455-01/806-5257/6je9h032b/index.html> (accessed on 9th February 2024).
- Eyng, R. (2019) *Redis: Pipelining, Transactions and Lua Scripts* [online]. Available at <https://rafaeleyng.github.io/redis-pipelining-transactions-and-lua-scripts> (accessed on 9th February 2024).
- Visual Paradigm. (2024) *What is Spike in Scrum?* [online]. Available at <https://www.visual-paradigm.com/scrum/what-is-scrum-spike/> (accessed on 9th February 2024).
- Ugarte, A. (2024) *What Is Thread-Safety and How to Achieve It?* [online]. Available at <https://www.baeldung.com/java-thread-safety> (accessed on 9th February 2024).
- Graefe, G. (2016) *Revisiting optimistic and pessimistic concurrency control* [online]. Available at <https://www.labs.hpe.com/techreports/2016/HPE-2016-47.pdf> (accessed on 9th February 2024).
- Thornton, G. (2001) *Optimistic Locking with Concurrency in Oracle* [online]. Available at <https://www.orafaq.com/papers/locking.pdf> (accessed on 9th February 2024).
- Weston, T. (2011) *Optimistic and Pessimistic Concurrency Control with Shared Memory Models* [online]. Available at <https://baddotrobot.com/resources/concurrency-control-1.0-draft.pdf> (accessed on 9th February 2024).
- Apache Software Corporation. (2013) *Deadlocks* [online]. Available at <https://db.apache.org/derby/docs/10.0/manuals/develop/develop75.html> (accessed on 9th February 2024).
- Heil, C. (2022) *Is Kanban the better Agile Approach in a Highly-Regulated Environment?* [online]. Available at https://www2.deloitte.com/content/dam/Deloitte/de/Documents/risk/Deloitte_Agile_Approach_in_Highly_Regulated_Environments.pdf (accessed on 12th February 2024).
- Rehkopf, M. (2024) *Kanban vs. scrum: which agile are you?* [online]. Available at <https://www.atlassian.com/agile/kanban/kanban-vs-scrum> (accessed on 12th February 2024).
- Mauvius Group Inc. (2021) *The official guide to the kanban method* [online]. Available at https://resources.kanban.university/wp-content/uploads/2021/06/The-Official-Kanban-Guide_A4.pdf (accessed on 12th February 2024).
- Jenkins. (2024) *Pipeline Syntax* [online]. Available at <https://www.jenkins.io/doc/book/pipeline/syntax/> (accessed on 12th February 2024).

- Rodriguez, P. Haghishatkhah, A. Lwakatare, L, E. Teppola, S. Suomalainen, T. Eskeli, J. Karvonen, T. Kuvala, P. Verner, J, M. Oivo, M. (2017) 'Continuous Deployment of Software Intensive Products and Services: A Systematic Mapping Study', Journal of Systems and Software, Vol. 123, pp. 263-291. Available at: <https://doi.org/10.1016/j.jss.2015.12.015>
- Department of Defence. (2021) *DoD Enterprise DevSecOps Fundamentals* [online]. Available at <https://dodcio.defense.gov/Portals/0/Documents/Library/DoDEnterpriseDevSecOpsFundamentals.pdf> (accessed on 14th February 2024).
- NSA. (2023) *Defending Continuous Integration/Continuous Delivery (CI/CD) Environments* [online]. Available at https://media.defense.gov/2023/Jun/28/2003249466/-1/-1/0/CSI_DEFENDING_CI_CD_ENVIRONMENTS.PDF (accessed on 14th February 2024).
- OWASP Foundation. (2023) *OWASP Top 10 CI/CD Security Risks* [online]. Available at <https://owasp.org/www-project-top-10-ci-cd-security-risks/> (accessed on 14th February 2024).
- Wikström, A. (2019) *Benefits and challenges of Continuous Integration and Delivery - A Case Study* [online]. Available at <https://core.ac.uk/download/pdf/226768285.pdf> (accessed on 14th February 2024).
- Fairbanks, J. Tharigonda, A. Eisty, N, U. (2023) *Analyzing the Effects of CI/CD on Open Source Repositories in GitHub and GitLab* [online]. Available at <https://arxiv.org/pdf/2303.16393.pdf> (accessed on 14th February 2024).
- GoRetro. (2023) *Ways of Working* [online]. Available at <https://www.goretro.ai/glossary/ways-of-working> (accessed on 14th February 2024).
- Banerjee, A. (2017) 'System Changeover Policy and Precautions', International Journal of Advanced Engineering and Management, Vol. 2, No. 9, pp. 214-218. Available at: <https://doi.org/10.24999/IJOAEM/02090048>
- Smyth, D. (2020) *Changeover Techniques* [online]. Available at <https://smallbusiness.chron.com/changeover-techniques-34890.html> (accessed on 15th February 2024).
- IoRedis. (2024) *README - ioredis* [online]. Available at <https://ioredis.readthedocs.io/en/stable/README/#transparent-key-prefixing> (accessed on 15th February 2024).
- Rustagi, V. (2023) *Strategies To Run Old & new Systems Simultaneously Using The Same Database* [online]. Available at <https://metaorangdigital.com/blog/strategies-to-run-old-new-systems-simultaneously-using-the-same-database/> (accessed on 15th February 2024).
- Bigelow, J, S. (2020) *What are the types of requirements in software engineering?* [online]. Available at <https://www.techtarget.com/searchsoftwarequality/answer/What-are-requirements-types> (accessed on 16th February 2024).

- Karolis, N, Saulius, G. (2023) 'Causal Knowledge Modelling for Agile Development of Enterprise Application Systems', *Informatica* (Netherlands), Vol. 34 Issue 1, p 124. Available at: <https://doi.org/10.15388/23-INFOR510>
- Eduardo, M. (2022) 'Moscow Rules: A Quantitative Exposé', Lecture Notes in Business Information Processing, Vol. 445, pp 19-34. Available at: https://doi.org/10.1007/978-3-031-08169-9_2
- Kim, M. Park, S. Sugumaran, V. Yang, H. (2007) 'Managing requirements conflicts in software product lines: A goal and scenario based approach', *Data & Knowledge Engineering*, Vol. 61 Issue 3, pp 417-432. Available at: <https://doi.org/10.1016/j.datak.2006.06.009>
- Kanungo, S, Morena, D, R. (2023) 'Concurrency versus consistency in NoSQL databases', *Journal of Autonomous Intelligence*, Vol. 7 Issue 3. Available at: <https://doi.org/10.32629/jai.v7i3.936>
- Sanders, R. (1987), 'THE PARETO PRINCIPLE: ITS USE AND ABUSE', *Journal of Services Marketing*, Vol. 1 No. 2, pp. 37-40. Available at: <https://doi.org/10.1108/eb024706>
- Hashimi, A, L. Gravell, A. (2020), 'Spikes in Agile Software Development: An Empirical Study', 2020 International Conference on Computational Science and Computational Intelligence (CSCI). Available at: <https://doi.org/10.1109/CSCI51800.2020.00319>
- Hashimi, A, L. Abduldaem, A. Gravell, A. (2022), 'Common Spikes Success Factors: An Industrial Investigation within Agile Software Development', 2022 12th International Conference on Software Technology and Engineering (ICSTE). Available at: <https://doi.org/10.1109/ICSTE57415.2022.00008>
- Hashimi, A, L. Gravell, A. (2019) 'A Critical Review of the Use of Spikes in Agile Software Development', The Fourteenth International Conference on Software Engineering Advances. Available at: https://www.researchgate.net/publication/340446921_A_Critical_Review_of_the_Use_of_Spikes_in_Agile_Software_Development
- Kruchten, P. Nord, L, R. Ozkaya, I. Visser, J. (2012) 'Technical Debt in Software Development: from Metaphor to Theory Report on the Third International Workshop on Managing Technical Debt', ACM SIGSOFT Software Engineering Notes Advances, Vol. 37, No. 5, p. 36. Available at: <https://doi.org/10.1145/2347696.2347698>
- Sathre, J, Zambreno, J. (2008) 'Automated software attack recovery using roll-back and huddle', Des Autom Embed Syst, Vol. 12, pp. 243-260. Available at: <https://doi.org/10.1007/s10617-008-9020-4>
- Burghate, N. (2018) 'Work Breakdown Structure: Simplifying Project Management', *International Journal of Commerce and Management Studies (IJCAMS)*, Vol. 3, No. 2. Available at: <https://www.ijcams.com/wp-content/uploads/2018/06/WBS.pdf>

- Glas, T. Hedén, F. (2021) *Managing Technical Debt in XP Teams* [online]. Available at <https://fileadmin.cs.lth.se/cs/Education/EDAN80/Reports/2021/GlasHeden.pdf> (accessed on 19th February 2024).
- Microsoft. (2021) *Technical Spike* [online]. Available at <https://microsoft.github.io/code-with-engineering-playbook/design/design-reviews/recipes/technical-spike/> (accessed on 19th February 2024).
- Martins, J. (2023) *7 common causes of scope creep, and how to avoid them* [online]. Available at <https://asana.com/resources/what-is-scope-creep> (accessed on 19th February 2024).
- Raeburn, A. (2024) *The work breakdown structure (WBS) for project management: What it is and how to use it* [online]. Available at <https://asana.com/resources/work-breakdown-structure> (accessed on 19th February 2024).
- The Open Group. (2016) *Implementation and Migration Elements* [online]. Available at https://pubs.opengroup.org/architecture/archimate30-doc/chap13.html#_Toc451758082 (accessed on 19th February 2024).
- Jonkers, H. Proper, E. Lankhorst, M, M. Quartel, A, C, D. Iacob, M (2011) 'ArchiMate(R) for Integrated Modelling Throughout the Architecture Development and Implementation Cycle', 2011 IEEE 13th Conference on Commerce and Enterprise Computing. Available at: <https://doi.org/10.1109/CEC.2011.52>
- Silveria, A. (2020) *GraphQL Live Querying with DynamoDB* [online]. Available at <https://arxiv.org/pdf/2008.00129.pdf> (accessed on 21st February 2024).
- Lloyd, K. (2023) *iPlayer Schedules - BBC Partner Schedule Generator - Delta Update Technical Design*. BBC. Unpublished.

10 Appendix

10.1 Appendix A - Partnerships Objectives

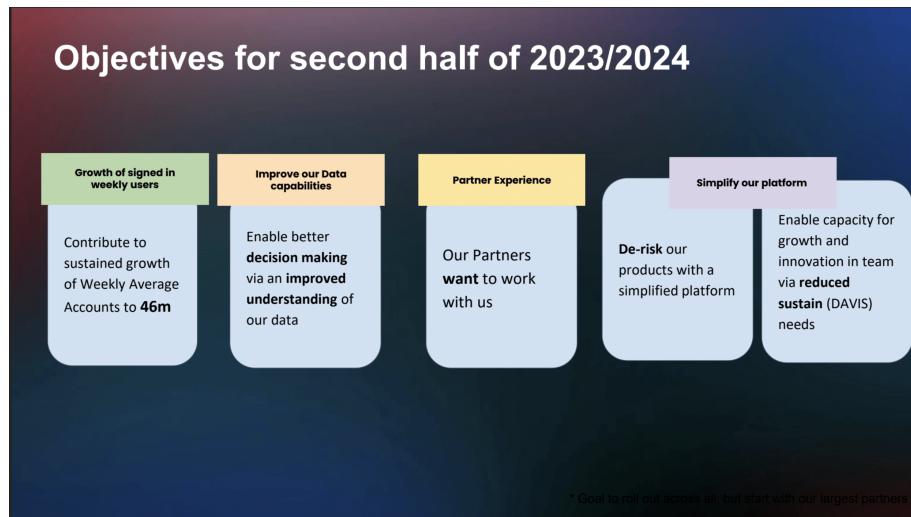


Figure 31: Image taken from a presentation given at a partnerships context setting event (BBC Partnerships, 2023).

10.2 Appendix B - Full schedules design including future external notification work

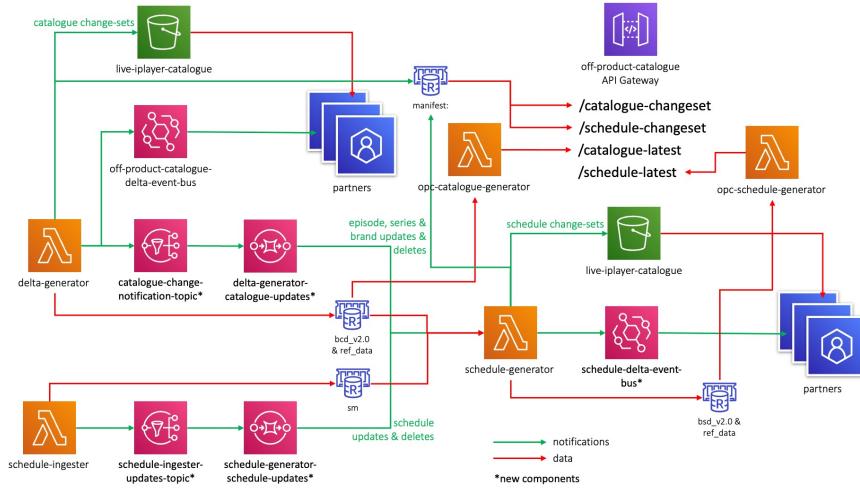


Figure 32: Full diagram of design for schedules pipeline, including future notifications to partners work (Lloyd, 2023).

10.3 Appendix C - Initial flow diagrams for how events would be processed

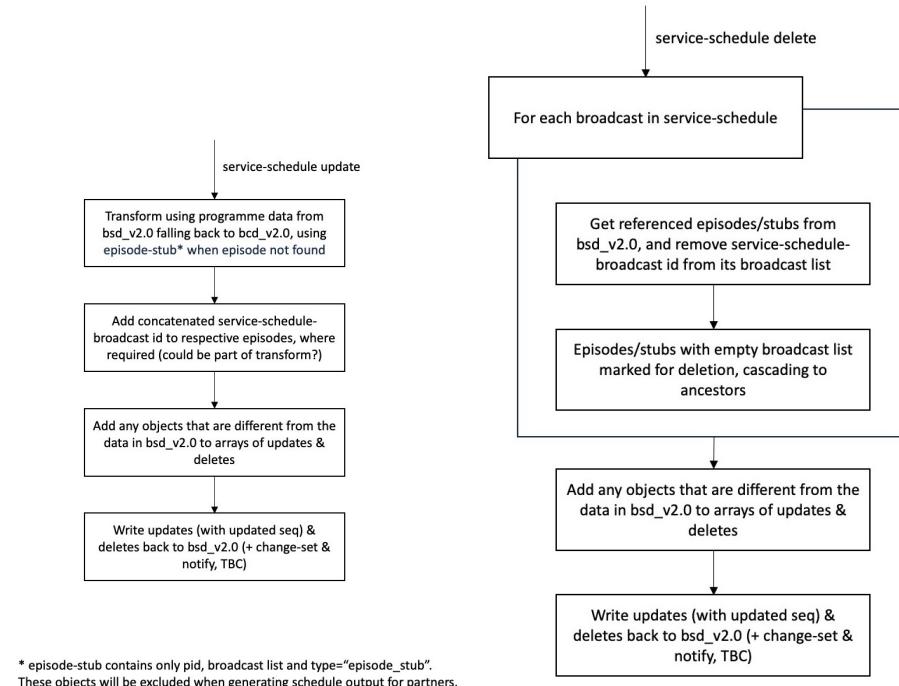


Figure 33: Flow diagrams for schedule events (Lloyd, 2023).

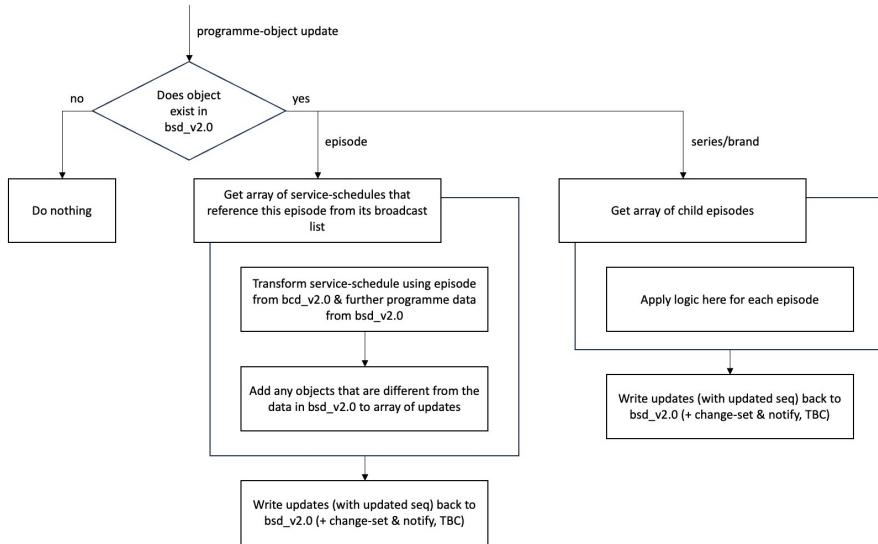


Figure 34: Flow diagram for catalogue/programme events (Lloyd, 2023).

10.4 Appendix D - Full ways of working diagram

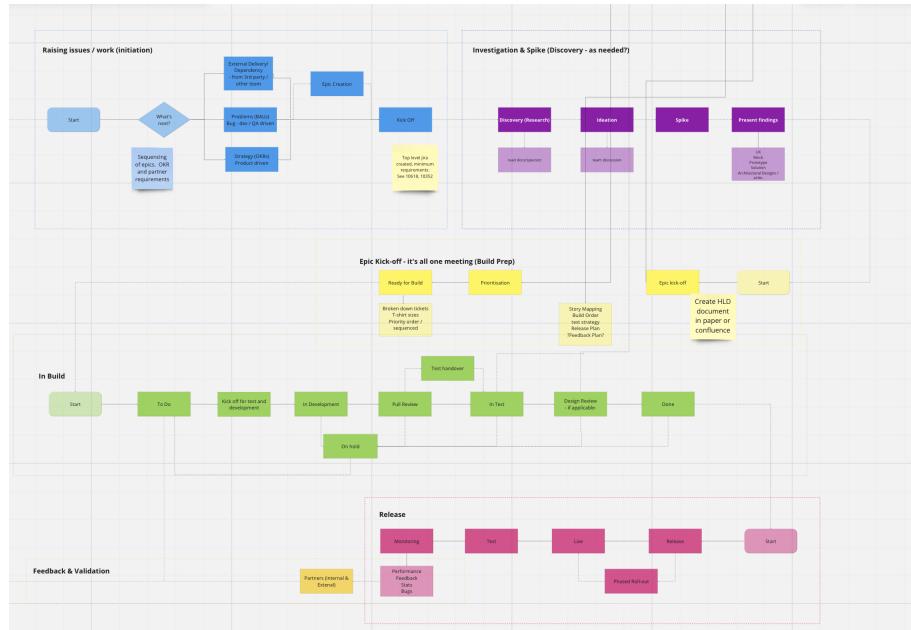


Figure 35: Full ways of working flow diagram used by SpaceChimp.

10.5 Appendix E - Full software spike document

Updating schedule-generator with deltas

Spike: Schedule-generator Deltas

- **Conducted by:** Fabrizio Colucci & Oli Bowker (with exec. consultants Ray Cheung & Karl Lloyd)
- **Backlog Work Item:** [IPLAYERTVV1-14412](#)

Goals

1. Validate proposed algorithms for dealing with service-schedule and catalogue updates and deletes flowing into the schedule-generator - no optimisations at this point
2. Outline what updates would be required to the existing cold-start mode to be compatible with this
3. Suggest what potential optimisations could be made, taking into consideration the current typical delta throughput from both sources

Method

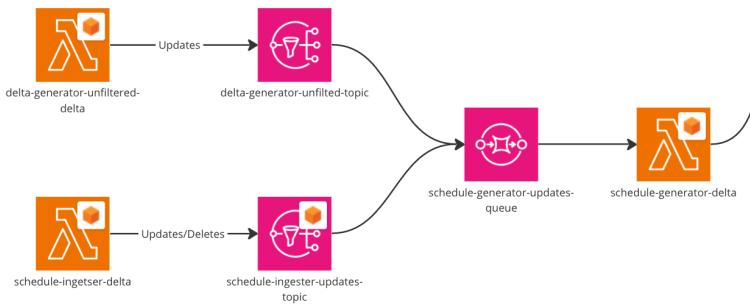
1. Build prototype to Int environment.
2. Use unit tests to cover list of scenarios in the above parent ticket.
3. Promote spike to Test environment to analyse capability of handling throughput of real data.
 - a. Add a delta-generator outgoing SNS for notifications. Send only on Int & Test?
 - b. Other infrastructure exists in schedule-generator spike branch - attach incoming SQS to existing schedule-ingester SNS, etc. Deploy this in parallel with current schedule-generator.
4. Analyse the complexity/maintainability of solution.

Test scenarios

See [IPLAYERTVV1-13966](#) description for initial list of scenarios to cover.

Evidence

The [repo](#) and unit tests for algorithm and subsequent AWS architecture. The delta lambda can be found [here](#) and the logs for the lambda can be found [here](#).



New architecture for the schedule-generator

Metrics have been setup in [grafana](#) to show throughput on TEST, under the Schedule Generator (SPIKE work) tab.

- Metric to track number of updates and deletes per type (episode/series/brand)
- Metric to track number of ignored updates and delete per type (episode/ancestor)
- Metric to track general lambda errors
- Metric tracking average run time, invocations and oldest message in queue.

Conclusions

Approach

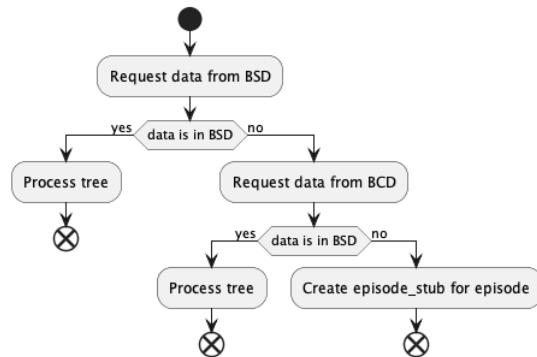
The final approach consisted of repositories requiring change.

- A spike [repo](#) was created that housed the new generator code. This code writes to the redis keyspace `{bsd_v2.0_spike}` and `{bsd_v2.0_ref_data_spike}`. Unit tests were written that covered scenarios specified in the initial [brief](#).
- Updates to the delta generator were added to add an SNS topic on the unfiltered lambda, which posts **UPDATES** events only to the schedule generator SQS.

On schedule **UPDATE** all broadcasts are iterated through and respective data episode/series/brand data is persisted into the `{bsd_v2.0_spike}` keyspace by the below algorithm. With broadcast lists of episodes being populated with an array of

`sid/broadcast_pid` identifiers. This list is used to know what schedules need updating when episode/series/brand updates come in.

- Edge case - broadcast_pid or episode is changed in an existing schedule.
Remove old identifier from the old referenced episode, if that broadcast_list is empty then remove the episode.



Schedule update tree building logic.

On schedule **DELETE** all broadcast are iterated through, with associated `sid/broadcast_pid` identifiers being removed from the episodes broadcast_list. If this list ends up being empty, the episode is removed from {bcd_v2.0_spike} alongside the schedule.

On episode **UPDATE** get the new episode data from {bcd_v2.0} and then append the broadcast_list and new seq.

On series/brand **UPDATE** get the new series/brand from {bcd_v2.0} and then append a new seq number to it. Then determine the linked episodes in {bcd_v2.0_spike} and update the relevant schedules in their broadcast_lists.

Catalogue **DELETES** are not explicitly handled.

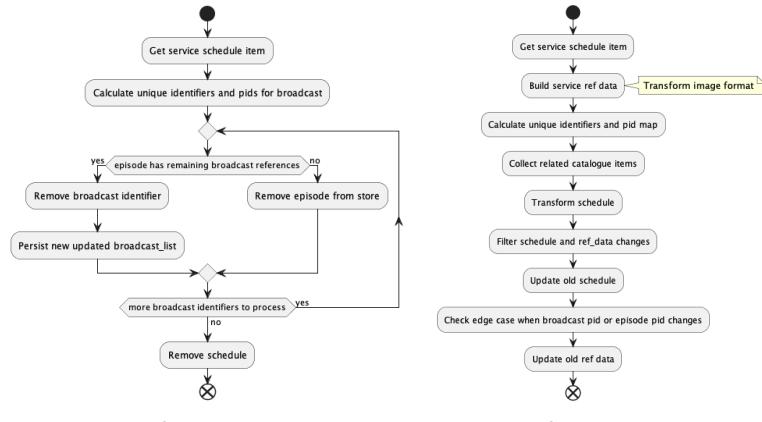
Coldstart

Coldstart functionality was not fully implemented during the spike however below is a list of things that would need to be handled.

- Standard ‘cut-off’ for the lambda SQS mappings would need to be added for the SQS going into the schedule generator. This also includes perging the queue.
- Protection when the delta-generator coldstarts. Implement some form of threshold of amont of updates that can be sent to the schedule-generator at once. If this threshold is breached, coldstart the schedule generator as we do in the aggricat.
- Purge all catalogue objects from `{bsd_v2.0}` and compile episodes/series/brands.

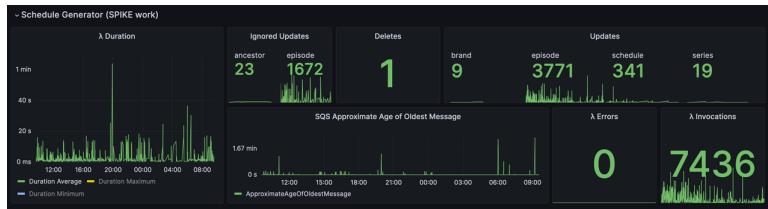
Garbage Collector

The current system only handles the deletion of episodes and schedules from the `{bsd_v2.0_spike}` keyspace. It was determined a single job ran daily to clean up the series and brands would be better as all data from the `{bsd_v2.0}` keyspace would have to be read and parsed for every delete. Considering deletes only happen once a day, we may as well tidy up once a day.



Monitoring

Below is moniroing stats for a 24 hour period. It's important to note the ignored updates, actual updates and deletes will all be higher when we do this for real. This is due to the spike not coldstarting, therefore there are less schedules in which in the store. However this does nicely show a *maximum* amount of events we're likely to recieve.



Pros	Cons
Schedules can somewhat rely on its own data store after initial setup of data is done.	Large duplication of redis data across {bsd_v2.0} and {bcd_v2.0}.
New code and infrastructure is all things that we have done before, nothing new to learn.	Hard to parallelise if we choose to.
Changesets should be easier to implement using a solution like this.	Catalogue and Schedule pipelines are not separate, (however this is already how it works on LIVE).

Potential Improvements

- Schedule-generator
 - Parallelise schedule updates triggered by episode/series/brand changes. An episode change then triggers updates for any schedule that it refers to, brand and series also do this but for the entire tree. An episode/series/brand present in BBC One can therefore will spin off a schedule change for all regional and hd variants. If a brand like eastenders was to change for whatever reason, this would result in 100s of schedule changes. This is a must.
 - Remove all catalogue deletes from schedule generator, and tidy up anything with an empty broadcast_list, and its parents in the garbage collector. This would simplify the current schedule-generator code and make garbage collection easier.
 - Implement episode→broadcast mappings in DynamoDB instead of having broadcast_list in episodes. This would remove the need for redis duplication and would allow us to parallelise using optimistic locking in the future if we choose/need to.

- This could be done at the ingestor, once more simplifying the schedule-generator even more.
- Delta-generator
 - More filtering could be applied as we only care if titling, synopses, subtitles, (images?) or warning_text has been updated. If a dynamoDB mapping approach is taken, then a lookup could be done and only relevant pids would be sent on top of this (maybe over engineering).
 - Make code more explicit that it wants to publish to SNS, don't rely on their not being a manifest key and the version not being v1.1.

Next Steps

If we are happy to continue with the above approach, Oli and Fab will create/slice tickets and a kick off can be planned.